

A voting system case study: Pret-A-Voter

Panagiotis Grontas

$\mu\Pi\lambda\forall$ - CoReLab Crypto Group

26/06/2013

Verifiable Elections With Receipts I

Cryptographic Voting Receipts

- Assure Voters that their vote is counted
- While maintaining ballot secrecy and avoiding coercion

Advantages

- Mathematically ensure election integrity
- Allow software independence [Riv08]
 - Avoid trust on proprietary hardware
 - Detect errors early (at the polling station)
 - Rerun the elections without software

Verifiable Elections With Receipts II

Software independence

A voting system is software independent if an undetected change or error in its software cannot cause an undetectable change or error in the election outcome.

How:

- Paper trail
- Cryptography

Voting with receipts

- Input choices on the touch screen of the voting machine
- Printer prints the receipt
- Voter checks the printout
- Accept or Change Vote
- Mechanism for ballot secrecy
- Receipt
- Mechanism for receipt secrecy
- Anonymization
- Tallying
- Verification: Web Bulletin Board to enter receipt SN and verify that it is included in the tally


Visual Cryptography

- Encrypt written material in a **perfectly secure** way
- Decryption is done by human visual system
- Ciphertext and secret key are two pages of papers
- Ciphertext can be transmitted by fax or published
- Secret key is printed on transparent paper
- Each one on its own is indistinguishable from random noise
- When placed on top of each other we have decryption
- A visual one time pad
 - Each page of cipher text is decrypted by its own transparency
- Can be extended to k of n threshold secret sharing scheme [NS95]


Receipts with Visual Cryptography - [Cha04]

- Receipts consist of two layers laminated together
- The printer prints the selection of the voter but does not conclude the job
- Voter selects to keep top or bottom layer and then the printing concludes
- Machine keeps the same layer in digital form
- Neither layer is readable on its own, but both encode the vote
- The other layer is destroyed by the voting personnel and the voting booth (securely?)


Encoding the vote with visual cryptography I




Encoding the vote with visual cryptography II




Newspaper



Top layer



Bottom layer



Laminated

PretAVoter - Introduction

- End To End Verifiable Voting System
- A design of ballots for electronic voting
- Creator: Peter Ryan (University of Luxemburg) (2004)
- Based on a previous cryptographic scheme by Chaum (2004)
- Influenced a subsequent cryptographic scheme by Chaum
- First implementation by University of Surrey (2007)
- Still in active development

PretAVoter - Characteristics

- Integrity
- Voter Privacy
- Receipt Freeness
- Voter Verifiability
- Voter Friendly
- Modular: Can implement many cryptographic voting protocols (encryption, zero knowledge etc.)
- Versatile: Can be used with many social choice functions
- Requires supervised voting infrastructure

Assumptions

- Voter registration and authentication is outside the scope of the system
- The ballot forms are properly managed (remain confidential, cannot be forged) between their creation and used
- Voting occurs in a voting booth with the traditional privacy characteristics
- There exists a tamper - free public bulletin board that can be used for verification

The PretAVoter Ballot I

Main idea

The vote is recorded on a different frame of reference for each ballot.

As a result the system **encrypts the ballot not the vote**

- The ballot consists of two columns: one for the candidate list and one for the voter selection.
- The candidate list is random for each ballot
- The voter marks or ranks the candidates on his/her ballot
- The candidate list **must** be disposed
- The actual vote is the mark or rank that is recorded on the vote column
- No encryption is needed, since the vote is useless without the candidate list
- The vote column is used as a receipt for voter verifiability

The PretAVoter Ballot II

- Since the vote order is different, a source of voter bias is removed but compliance might be affected

Candidates	Your vote
Obelix	
Idefix	
Asterix	
Panoramix	
	<i>7rJ94K</i>
Destroy	Retain

Your Vote
X
<i>7rJ94K</i>
Retain

Candidates	Your vote
Asterix	
Idefix	
Panoramix	
Obelix	
	<i>N5077t3</i>
Destroy	Retain


- There is a base ordering of a candidates
- The ordering on each ballot is a permutation of the base

The PretAVoter Ballot III

- The difference (offset) is encrypted and printed on each ballot
- The encrypted value is called onion, referring to the processing of the ballot in the first version of the systems
- The voting equipment records the index/ranking and transmits it along with the encrypted onion
- Only the voter knows the order of the candidates that he voted on
- **Extra bonus:** side channel attacks are side stepped
 - No encryption takes place at the booth (the onion is created in advance)
 - The vote without the candidates is useless

Vote Processing I

- 1 Mixing
- 2 Decrypting
- 3 Tallying




Decryption Mixnet I

- n voters
- m candidates
- Each mix server performs 2 mixes
 - k mix servers
 - $2k$ key pairs
- Ballot generation - Layered Generation
 - $2k$ random values $\{g_{i,j}\}_{0,0}^{n-1,2k-1}$ - germs - for each voter
 - the germs will be used for
 - **randomised encryption**
 - Onion for ballot i : $D_{i,2k} = \text{Enc}(g_{i,2k-1}, \text{Enc}(\dots \text{Enc}(g_{i1}, \text{Enc}(g_{i0}, D_{i0})), \dots))$
where D_{i0} random value for initialisation
 - **offset calculation**
 - $\theta_i = \sum_{j=0}^{2k-1} d_j \mod m$ where $d_j = \text{hash}(g_{ij})$

Decryption Mixnet II

- Vote i is a pair $(r_i, D_{i,2k})$ where r_i is the index of the marked candidate
- Mix server operation
 - Input $L_{2j} = \{r_i, D_{i,2j}\}_{i=0, j=k-1}^{n-1, 0}$
 - Decrypt the onion with the private key (g_{ij}, D_{ij})
 - Apply the hash to the germ value of the onion $d_{ij} = \text{hash}(g_{ij})$
 - Calculate new shift $r_{ij} = r_{i,j-1} - d_{ij}$
 - Sort the list by the value of the onions
 - Audit with RPC

Decryption Mixnet III



Threats

- Confidentiality of ballots, before elections. There is a need to trust the authority and the intermediates
- Chain Voting: Coercion by smuggling a ballot, pre-voting, and require a replacement ballot
- Kleptographic Channels: Choose special values of cryptographic variables in order to leak information to a colluding party
 - Tagging attack in mixnets

Common Theme: Too much trust in a single authority

Solution: Distributed Generation of Ballots

Distributed Generation Of Ballots I

- The ballot forms will be generated by l entities called *clerks*
- Each clerk contributes to the generation of a common encrypted seed
- The candidate list is constructed from the seed
- To reveal the seed values all clerks must be corrupted
- Encryption using exponential ElGamal (also works for RSA)
- Tellers conduct tallying using a reencryption mixnet

Distributed Generation Of Ballots II

Operation

- Generate 2 encrypted '*onions*':
- $onion_L$
 - will create the random candidate list
 - will be read (decrypted) by the voting machine in order to display the candidates to the voter
- $onion_R$
 - will be used for voter verifiability
 - and reconstruction of the candidate list

Distributed Generation Of Ballots III

Parameters

- p, q safe primes
- α, γ generators of \mathbb{Z}_q^*
- β_T the public key of the tellers
- β_R the public key of the vote recording machines

Clerk C_0

- Generate a batch of initial seeds s_i^0 - one for each ballot
- Create encryption with randomness x_i^0, y_i^0 for
- voting machine $E_R(\gamma^{-s_i^0}, x_i^0) = (\alpha^{x_i^0}, \beta_R^{x_i^0} \cdot \gamma^{-s_i^0})$
- teller $E_T(\gamma^{-s_i^0}, y_i^0) = (\alpha^{y_i^0}, \beta_T^{y_i^0} \cdot \gamma^{-s_i^0})$

Distributed Generation Of Ballots IV

Clerk C_l

- Receive batch of onions from previous clerk
- Generate fresh randomness
- Perform reencryption and shuffle

Final onions

- $onion_L = E_R(\gamma^{-\sum_i s_i}, \sum_i x_i)$
- $onion_R = E_T(\gamma^{-\sum_i s_i}, \sum_i y_i)$

Correct operation \Rightarrow the 'messages' match.

Voting with distributed ballots I

- The onions are stored in encrypted form and are presented to the voter on the ballot form

$onion_L$	$onion_R$

- Device reads and decrypts the left hand onion s is revealed
- The candidate permutation is reconstructed as a function of s
- The device should not see the right onion

Voting with distributed ballots II

Voting

- The voter now faces a classical ballot
- Marks candidate index
- Embed selected candidate index into seed by reencryption
 $vote = E_T(\gamma^{r - \sum_i s_i}, \sum_i y_i)$
- Mixing proceed as usual
- Decryption: plaintexts are of the form $\gamma^{r - \sum_i s_i}$
- Solve discrete logarithm
- Select seeds so that the exponent is not too large and the search space is bounded
- **Solution:** Use Paillier Cryptosystem

Overview of STV Voting I

Objective:

Reduce vote waste, for candidates that cannot win the election

- Each voter gets a single vote, but 'this vote' can be used by many candidates
- When a candidate is eliminated, 'his' votes are transferred
- Each voter ranks candidates
- Ranking selects the candidates to receive the votes in case favourites are eliminated
- **Tallying**
 - Partition ballots according to the first choice of the voter
 - The candidate with the fewest votes is eliminated
 - His votes are transferred according to the second choices
 - Repeat until there is only one candidate

Overview of STV Voting II

Remarks

- **Quotas:** The number of votes that assure election
- If a candidate has reached the quota, he cannot be eliminated
- But the surplus of votes has to be transferred, as well
- The voter should not provide a complete ranking of the candidates.
- A ballot with all the candidates eliminated is discarded

Overview of STV Voting III

Useful Characteristics

- **Unused Preferences:** Some of the preferences might not be used, since the winner might get declared earlier.
- **Transfer history** is irrelevant. We are not interested for the previous candidates since they are eliminated
- **Lazy evaluation** can be applied. We can work on the candidates one by one.

Difficulties

- **The italian attack** Many preferences introduce a covert channel, as lower ranked preferences can be used to identify the voter

Incorporating STV into PretAVoter [Hea07] I

Attempt 1

- Encode a partial ranking in the vote
- Keep the cyclic shift
- But this might leak too much information because it preserves the relations found in the canonical ordering
- An example:

People's Front of Judea 1	
Judean People's Front 1	
People's Front of Judea 2	
Judean People's Front 2	
Judean People's Front 3	
People's Front of Judea 3	

	1
	6
	3
	2
	4
	5

- We must embed a complete permutation

Incorporating STV into PretAVoter [Hea07] II

Attempt 2 - ElGamal

- Encode a permutation in the ElGamal exponent in a way that addition/multiplication yields permutation composition
- Cannot be done

Attempt 3 - RSA onions

- Each layer encodes a permutation that can be recovered and applied
- **Problems**
 - It is evident how many choices were made
 - This can be used to partition ballots and leak information
 - No random padding can be added, since it will affect the election results
 - The candidates that were not ranked can be used as a covert channel (Italian attack)

Incorporating STV into PretAVoter [Hea07] III

Solution: Multi valued reencryption mixes

- Encode the vote as a sequence of cryptograms $\{c_1, \dots, c_k\}$ that decrypts to a candidate index
- The order of the cryptograms reflects the ranking of the candidates by the voter, ie c_1 is the encryption of the index of the favourite candidate
- A dummy candidate 'STOP' id included that marks the end of the user ranking
- A complete ranking is given by using a random permutation of the remaining candidates
- **Lazy Tallying:** First decrypt the heads
- First preferences are counted
- If a candidate is eliminated then the head is appended at the end, an anonymising mixnet is run and the heads are again decrypted
- Stop when a winner (a candidate with more than half the votes) is found

STV with PretAVoter: An example I

Ballot Form

- Each ballot form contains a signed hash
- Hash serves as a key to lookup $k + 1$ onions
 - The first k correspond to the candidates $\{i, z\}_{PK_T}, i = 1 \dots k$ and are presented in random order
 - The last is the encryption of the stop candidate $\{0, z\}_{PK_T}$

Initial Ballot Form

Left Side: Registers

Right Side: Tellers

Encryption of candidate indexes

$[Alice]_{PK_R}$	$[Alice]_{PK_T}$
$[Bob]_{PK_R}$	$[Bob]_{PK_T}$
$[Charlie]_{PK_R}$	$[Charlie]_{PK_T}$
$[Duncan]_{PK_R}$	$[Duncan]_{PK_T}$
$[Esther]_{PK_R}$	$[Esther]_{PK_T}$
	$[STOP]_{PK_T}$

STV with PretAVoter: An example II

- Distributed Ballot Generation

$[\text{Charlie}]_{PK_R}$	$[\text{Charlie}]_{PK_T}$
$[\text{Bob}]_{PK_R}$	$[\text{Bob}]_{PK_T}$
$[\text{Esther}]_{PK_R}$	$[\text{Esther}]_{PK_T}$
$[\text{Alice}]_{PK_R}$	$[\text{Alice}]_{PK_T}$
$[\text{Duncan}]_{PK_R}$	$[\text{Duncan}]_{PK_T}$
	$[\text{STOP}]_{PK_T}$

Reordered Ballot Form

Dummy Position Candidate Unchanged

Only the registrars can read the candidate list on the left side

Only the tellers can read the onions on the right side

- Each ballot can be printed by a subset of registrars after being retrieved from a bulletin' board

STV with PretAVoter: An example III

Voting

Voter sees the decrypted ballot form, without the 'STOP' candidate and ranks the

Charlie	
Bob	3
Esther	1
Alice	
Duncan	2
STOP	
	w4JKz

candidates (partially)
hash to the bulletin board

Booth sends the ranking and the

STV with PretAVoter: An example IV

PostVoting Bulletin Board Actions

- First Available ranking is assigned to 'STOP' candidate
- Rest of empty voters are filled from top to bottom
- Use the hash to retrieve candidate list
- Sort Based on ranking

5
3
1
6
2
4
w4JKz

$[\text{Charlie}]_{PK_T}$
$[\text{Bob}]_{PK_T}$
$[\text{Esther}]_{PK_T}$
$[\text{Alice}]_{PK_T}$
$[\text{Duncan}]_{PK_T}$
$[\text{STOP}]_{PK_T}$
w4JKz

$[\text{Esther}]_{PK_T}$
$[\text{Duncan}]_{PK_T}$
$[\text{Bob}]_{PK_T}$
$[\text{STOP}]_{PK_T}$
$[\text{Charlie}]_{PK_T}$
$[\text{Alice}]_{PK_T}$
w4JKz

STV with PretAVoter: An example V

Anonymising and tallying

- Apply STV
- Decrypt candidates one at a time, until stop is reached
- After each decryption, another round of anonymisation occurs

Esther	Duncan	STOP
$[\text{Duncan}]_{PK_T}$	$[\text{Bob}]_{PK_T}$	$[\text{Charlie}]_{PK_T}$
$[\text{Bob}]_{PK_T}$	$[\text{STOP}]_{PK_T}$	$[\text{Alice}]_{PK_T}$
$[\text{STOP}]_{PK_T}$	$[\text{Charlie}]_{PK_T}$	$[\text{Esther}]_{PK_T}$
$[\text{Charlie}]_{PK_T}$	$[\text{Alice}]_{PK_T}$	$[\text{Duncan}]_{PK_T}$
$[\text{Alice}]_{PK_T}$	$[\text{Esther}]_{PK_T}$	$[\text{Bob}]_{PK_T}$

Condorcet Voting

- The voter provides a full ranking of the candidates
- All candidates engage in pairwise comparisons (duels)
- The winner of each duel is the candidate that is preferred by the majority of the voters

An example - duel between the candidates c_A and c_B

For each ballot b do

if c_A is ranked higher than c_B in b then

$$w_A \leftarrow w_A + 1$$

else

$$w_B \leftarrow w_B + 1$$

duel - winner \leftarrow candidate with $w = \max(w_A, w_B)$

- The winner of the election is the one that wins **all** the duels
- There might not be a winner, so a completion method is used (but this is a backend issue)

Ballots for Condorcet Voting [CM05]

- For c candidates we have $\frac{c(c-1)}{2} \rightarrow O(c^2)$ ballots
- Each ballot is a *yes/no* choice for candidates i, j
- A yes choice indicates preference for candidate i over candidate j
- Each yes/no ballot has its own onion $\langle i, j \rangle$
- Layered construction of onions for use in decryption mixnet
- Each voter submit $c \times c$ votes with the preferences and the onions
- The vote is run through a decryption mixnet for anonymisation
- **Remark:** There can be no correlation that all c^2 votes belong to the same voter
- A summation matrix is used to declare the winner

References I



David Chaum.

Secret-ballot receipts: True voter-verifiable elections.
Security & Privacy, IEEE, 2(1):38–47, 2004.



Michael R. Clarkson and Andrew C. Myers.

Coercion-resistant remote voting using decryption mixes.
In In Frontiers in Electronic Elections (FEE 2005), 2005.



David Chaum, Peter Y. A. Ryan, and Steve A. Schneider.

A practical voter-verifiable election scheme.
In ESORICS, pages 118–139, 2005.



James Heather.

Implementing stv securely in pret a voter.
In Computer Security Foundations Symposium, 2007. CSF'07. 20th IEEE, pages 157–169. IEEE, 2007.



Moni Naor and Adi Shamir.

Visual cryptography.
In Advances in Cryptology EUROCRYPT'94, pages 1–12. Springer, 1995.



Peter YA Ryan, David Bismark, JA Heather, Steve A Schneider, and Zhe Xia.

The prêt à voter verifiable election system.
IEEE transactions on information forensics and security, 4(4):662–673, 2009.



Ronald L Rivest.

On the notion of “software independence” in voting systems.
Philosophical Transactions of the Royal Society A: Mathematical, Physical and Engineering Sciences, 366(1881):3759–3767, 2008.

References II



Peter Y. A. Ryan and Steve A. Schneider.

Prêt à voter with re-encryption mixes.

In *ESORICS*, pages 313–326, 2006.



Peter YA Ryan.

Prêt à voter with paillier encryption.

Mathematical and Computer Modelling, 48(9):1646–1662, 2008.



Zhe Xia, Chris Culnane, James Heather, Hugo Jonker, Peter YA Ryan, Steve Schneider, and Sriramkrishnan Srinivasan.

Versatile prêt à voter: Handling multiple election methods with a unified interface.

In *Progress in Cryptology-INDOCRYPT 2010*, pages 98–114. Springer, 2010.